

Six Weeks Industrial Training Project Report on

**SKIN CONDITION DETECTION USING FACE RECOGNITION**

SUBMITTED IN PARTIAL FULFILLMENT OF THE REQUIREMENT FOR  
THE AWARD

OF THE DEGREE OF


**BACHELOR OF TECHNOLOGY**

IN

**COMPUTER SCIENCE & ARTIFICIAL INTELLIGENCE**

**Batch**

**(2022-2026)**



**SUBMITTED BY**

Neeraj Thakur  
12200981

**SUPERVISED BY**

Ms. Sangeeta Bhatti

**DEPARTMENT OF COMPUTER SCIENCE AND  
ENGINEERING DAV UNIVERSITY JALANDHAR-PUNJAB 144012  
JALANDHAR -144001, PUNJAB**

## **ACKNOWLEDGMENT**

I would like to express my sincere gratitude to my project guide, Ms. Sangeeta Bhatti, for their valuable guidance, constructive feedback, and continuous support throughout the department of computer science Engineering for providing the necessary resources and an encouraging academic environment.

My appreciation extends to all faculty members and classmates whose suggestions and cooperation helped strengthen the quality of this work. Lastly, I am deeply grateful to my family for their constant encouragement and support, which motivated me to complete this project successfully.

## **DECLARATION**

I, **Neeraj Thakur**, hereby declare that the work which is being presented in this project titled “**Skin Condition Detection using Face Recognition**” by me, in partial fulfilment of the requirements for the award of Bachelor of Technology (B. Tech) Degree in “Computer Science and Artificial Intelligence” is an authentic record of my own work carried out under the guidance of **Ms. Sangeeta Bhatti**. To the best of my knowledge, the matter embodied in this report has not been submitted to any other University/ Institute for the award of any degree or diploma.

NEERAJ THAKUR  
12200981

## ABSTRACT

The integration of Artificial Intelligence (AI) into dermatological diagnostics represents a paradigm shift in personalized healthcare. This industrial training project report comprehensively details the design, development, and deployment of the "Skin Condition Prediction Module," a critical component of the "Nepika" mobile application. The primary objective of this training was to engineer a robust machine learning pipeline capable of analyzing facial imagery to detect and classify specific dermatological conditions, including acne, dark circles, and skin tone variations, thereby democratizing access to preliminary skin analysis.

The project was executed over a four-week intensive training period, adhering to the rigorous standards of the System Development Life Cycle (SDLC). The core technical challenge addressed was the high variability in user-generated content, specifically selfie images with diverse lighting, angles, and resolutions. To overcome this, the project utilized advanced Computer Vision techniques and Deep Learning architectures. Specifically, the report elucidates the implementation of Convolutional Neural Networks (CNNs), utilizing Transfer Learning methodologies with the VGG16 and ResNet50 architectures to achieve high-fidelity feature extraction.

Beyond the algorithmic implementation, this report provides an exhaustive examination of the theoretical frameworks underpinning modern AI, the data preprocessing pipelines required to normalize biological data, and the software engineering practices employed to wrap the intelligence model into a scalable RESTful API using Python and Flask. The outcomes of this project indicate that AI-driven analysis can achieve diagnostic accuracy comparable to preliminary visual inspections by professionals, validating the "Nepika" module as a viable tool for consumer-grade skincare guidance. This document serves as a complete record of the training experience, satisfying the academic requirements of DAV University while contributing to the body of knowledge in applied medical AI.

## INTERNSHIP COMPLETION CERTIFICATE

InStep Technologies Pvt Ltd

Date: 14<sup>th</sup> August 2025

### To Whom It May Concern

This is to certify that **Neeraj Thakur** has successfully completed the **Data Science & Machine Learning Training Program** at **InStep Technologies Pvt. Ltd.** from **14th June 2025** to **14th August 2025**.

#### **His training covered:**

- **Server-Side Development** – Node.js, REST API creation
- **Database Systems** – DBMS & RDBMS, PostgreSQL, MySQL, MS SQL, schema design, normalization
- **System Design & Documentation** – UML diagrams, Software Requirements Specification (SRS), project architecture
- **Data Visualization** – Power BI, dashboard creation, dynamic data integration from APIs & databases
- **Data Science & AI Tools** – Python (Pandas, NumPy, Seaborn, scikit-learn, FastAPI, TensorFlow, Keras) for data preprocessing, analytics, and deployment
- **Machine Learning & Deep Learning** – Supervised/unsupervised learning, CNN, transfer learning, time series forecasting, and recommendation systems



### **Performance Overview:**

Neeraj displayed exceptional commitment and an eagerness to learn. His punctuality and professional behavior set a high standard for his peers. He was market-ready by the end of the program, capable of handling real-world projects with minimal supervision.

### **Learning & Application:**

Throughout the program, Neeraj engaged in both theoretical learning and practical assignments. He applied his skills to create server-side applications, integrate databases, design system architectures, and develop interactive dashboards connected to live data sources.

### **Conclusion:**

Neeraj has not only acquired a wide range of technical skills but also developed the professional mindset required to succeed in the technology industry. We believe he will be a valuable asset to any organization he chooses to join in the future.

Sincerely,  
Team Human Resources  
6280928238



## TABLE OF CONTENT

<b>Sr. No.</b>	<b>Content</b>	<b>Page No</b>
1	<b>Chapter 1: Introduction</b>	8 - 10
	1.1 Introduction to Industrial Training Context	
	1.2 Profile of the Training Organization	
	1.3 The Intersection of AI and Dermatology	
	1.4 Technological Landscape and Trends	
2	<b>Chapter 2: Title of the Project</b>	10 - 12
	2.1 Project Overview: The Nepika Ecosystem	
	2.2 Problem Statement and Clinical Relevance	
	2.3 Scope of the Study and Limitations	
3	<b>Chapter 3: Objectives</b>	12- 14
	3.1 Primary Technical Objectives	
	3.2 Training Program Learning Outcomes	
	3.3 Task Assignment and Weekly Breakdown	
4	<b>Chapter 4: Steps to Achieve Objectives</b>	14 - 17
	4.1 System Development Life Cycle (SDLC) Methodology	
	4.2 Requirement Analysis and Hardware Specification	
	4.3 Data Collection, Ethics, and Augmentation Strategies	
	4.4 Theoretical Framework: Deep Learning and CNNs	
	4.5 Mathematics of Backpropagation and Optimization	
5	<b>Chapter 5: Coding and Implementation</b>	18 - 23
	5.1 Development Environment Setup (Anaconda/CUDA)	
	5.2 Algorithmic Design and Workflow	
	5.3 Model Architecture: Transfer Learning with VGG16	
	5.4 Detailed Code Implementation and Logic	
	5.5 API Development with Flask	
	5.6 Testing, Validation, and Performance Metrics	
6	<b>Chapter 6: Conclusions and Recommendations</b>	24 - 26
	6.1 Summary of Findings	
	6.2 Critical Analysis of Limitations	
	6.3 Future Scope: Real-time Analysis and GANs	
7	<b>Appendices</b>	26 ..
	Appendix A: Source Code Snippets	

# **Chapter 1: Introduction**

## **1.1 Introduction to Industrial Training Context**

The transition from academic learning to professional application is a critical phase in the development of a computer science engineer. The Bachelor of Technology curriculum provides a robust foundation in theory—covering the mathematical basis of computation, the logic of algorithms, and the syntax of programming languages. However, the industrial environment introduces complexities that cannot be simulated in a classroom: the unpredictability of real-world data, the constraints of deployment environments, and the rigor of collaborative software development.

This report documents the "Four Weeks Industrial Training" undertaken as a mandatory component of the B.Tech degree. The training was not merely a passive observation period but an active engagement in the development lifecycle of a commercial-grade product. The focus domain, Artificial Intelligence (AI) and Machine Learning (ML), was selected due to its transformative impact on the modern technological landscape. Specifically, the training focused on Computer Vision—the science of enabling machines to "see" and interpret visual data—applied to the healthcare sector.

The project developed, the "Nepika Skin Condition Prediction Module," serves as a quintessential example of modern software engineering, combining backend logic, data science, and cloud computing. This introduction sets the stage for the detailed technical analysis that follows, positioning the project within the broader context of engineering education and industry trends.<sup>1</sup>

## **1.2 Profile of the Training Organization**

The training was conducted at a forward-thinking technology firm specializing in digital health solutions. While the specific proprietary details of the company's internal operations are confidential, the organizational structure and technical culture were pivotal to the learning experience.

The organization operates on an Agile methodology, characterized by iterative development, continuous feedback, and cross-functional collaboration. The AI division, where this training took place, is tasked with researching and deploying state-of-the-art models to solve consumer problems. The infrastructure provided to trainees included access to cloud-based GPU clusters

(AWS EC2 instances with NVIDIA Tesla GPUs) and proprietary data repositories, enabling high-performance model training that would be impossible on standard consumer hardware.

The mentorship structure involved daily "stand-up" meetings where progress was tracked, and technical blockers were addressed. This environment fostered a culture of problem-solving and technical accountability. The exposure to version control systems like Git in a team setting, where code conflicts must be resolved and branches merged, was a significant learning outcome of this organizational placement.

---

*(Placeholder Description: A hierarchical chart showing the flow from Project Manager -> Lead Data Scientist -> AI Engineer -> Trainee (Me).)*

### 1.3 The Intersection of AI and Dermatology

Dermatology is a unique field within medicine because diagnosis is primarily visual. A dermatologist observes patterns—color, texture, shape, and distribution of lesions—to form a diagnosis. This visual nature makes dermatology a prime candidate for automation via Computer Vision.

Historically, computer-aided diagnosis relied on "hand-crafted features." Engineers would manually write code to detect edges or measure redness. These systems were brittle; they failed if the lighting changed or if the skin tone varied. The advent of Deep Learning, specifically Convolutional Neural Networks (CNNs), revolutionized this. CNNs do not require manual feature definition. Instead, they learn features directly from the data.

The "Nepika" project leverages this capability. By training a neural network on thousands of images of skin conditions, the system learns to recognize the subtle, non-linear patterns that distinguish acne from a heat rash, or dark circles from natural shadowing. This capability addresses a massive gap in the market: the need for accessible, instant, and private skin advice.

### 1.4 Technological Landscape and Trends

The development of the Nepika module sits at the convergence of three major technological trends:

1. **Mobile Computing Power:** Modern smartphones possess cameras with high resolution and processors capable of running lightweight AI models, making apps like Nepika feasible.
2. **Transfer Learning:** The democratization of AI is largely due to Transfer Learning, where

models trained by tech giants (Google, Facebook) on massive datasets are fine-tuned by smaller teams for specific tasks. This project heavily utilizes this trend.

3. **Telemedicine:** Post-2020, there has been a global surge in remote healthcare. AI triage tools that can screen patients before they see a doctor are becoming standard infrastructure.
- 

## Chapter 2: Title of the Project

### 2.1 Project Overview: The Nepika Ecosystem

The project is titled "**Nepika: AI-Based Skin Condition Prediction Module.**"

To understand the module, one must understand the ecosystem it inhabits. Nepika is designed as a holistic skincare companion application. The full application architecture includes a frontend (built in React Native/Flutter) that handles user interaction, a product recommendation engine, and a tracking system for skincare routines.

However, the **Intelligence Layer**—the focus of this training report—is the brain of the operation. This backend module receives a raw image file, processes it, analyzes it for specific biological markers, and returns a structured analysis.

#### Core Functionality:

- **Input:** A facial image (Selfie).
  - **Processing:** Face detection, alignment, normalization.
  - **Analysis:** Multi-class classification (Acne, Dark Circles, Wrinkles, Normal).
  - **Output:** JSON response with confidence scores (e.g., {"acne": 0.85, "dark\_circles": 0.12}).
- 

*(Placeholder Description: Diagram showing User -> Mobile App -> API Gateway -> Python Flask Server -> Deep Learning Model -> Database -> Response.)*

### 2.2 Problem Statement and Clinical Relevance

The skincare industry is characterized by information asymmetry. Consumers are bombarded with products claiming to solve various issues, but they often lack an accurate understanding of their own skin conditions. Misidentification is common; for example, a user might mistake

fungal acne for bacterial acne, leading to incorrect treatment.

Furthermore, professional dermatological care is often inaccessible due to:

1. **Cost:** Consultation fees can be prohibitive.
2. **Availability:** Wait times for specialists can be weeks or months.
3. **Geography:** Rural areas often lack specialized clinics.

The Technical Problem:

From a computer science perspective, the problem is Fine-Grained Image Classification. Unlike distinguishing a cat from a dog (where the subjects are structurally different), distinguishing different skin conditions is subtle. An acne cyst and a bug bite look very similar. The system must be robust enough to handle:

- **Occlusion:** Hair or glasses covering the face.
- **Lighting:** Yellow bathroom light vs. blue daylight.
- **Pose:** Side profiles vs. frontal views.

## 2.3 Scope of the Study and Limitations

The scope of this training project was strictly defined to ensure achievability within the four-week timeframe.

**Inclusions:**

- **Conditions:** The model was trained to detect **Acne** (inflammatory lesions), **Dark Circles** (periorbital hyperpigmentation), and **Skin Tone** (Fair, Medium, Dark).
- **Data Scope:** The dataset comprised approximately 4,000 images, split between public datasets and web-scraped data.
- **Deployment:** The deliverable was a local REST API endpoint.

**Exclusions:**

- **Medical Diagnosis:** The app is a *cosmetic* tool, not a medical device. It provides "predictions," not "diagnoses."
  - **Rare Conditions:** Complex conditions like Melanoma or Eczema were excluded due to the high risk of false negatives and the need for specialized medical datasets.
  - **Frontend Development:** The creation of the mobile app UI was outside the scope of this backend intelligence training.
-

## Chapter 3: Objectives

### 3.1 Primary Technical Objectives

The development of the Nepika Skin Prediction Module was guided by specific, measurable technical objectives. These metrics served as the benchmarks for success during the training period.

1. **Classification Accuracy:** The primary goal was to achieve a validation accuracy of **>85%** for the binary classification of acne vs. healthy skin. This threshold is generally considered the minimum viability for a consumer application to be useful without frustrating the user.
2. **Inference Latency:** In a mobile application context, speed is paramount. The objective was to optimize the model and API such that the total time from receiving the image to returning the JSON result was **under 2 seconds** on a standard server configuration.
3. **False Positive Rate Minimization:** In skincare, telling a user they have acne when they do not can be psychologically damaging and lead to unnecessary product purchases. Therefore, a secondary objective was to prioritize high Precision over Recall, ensuring that positive detections were highly likely to be correct.
4. **Architectural Modularity:** The code needed to be written in a modular fashion—separating data loading, preprocessing, model definition, and training loops—to allow for future team members to update individual components without breaking the system.

### 3.2 Training Program Learning Outcomes

Beyond the immediate project deliverables, the industrial training was designed to instill specific engineering competencies in the trainee. These learning outcomes align with the overarching educational goals of the DAV University curriculum.

- **Advanced Python Proficiency:** Moving beyond basic scripting to object-oriented programming in Python, utilizing advanced features like decorators for the API and generators for memory-efficient data loading.
- **Deep Learning Literacy:** Gaining a deep intuitive and mathematical understanding of neural networks, including the mechanics of Gradient Descent, Backpropagation, and the role of activation functions.
- **Data Engineering:** Learning the often-overlooked skill of data cleaning—handling corrupt files, normalizing disparate image formats, and balancing class distributions to prevent model bias.

- **Professional Toolchain Mastery:** Becoming proficient with the standard stack of an AI engineer: Anaconda for environment management, Jupyter for prototyping, VS Code for production coding, and Git for version control.

### 3.3 Task Assignment and Weekly Breakdown

The project management followed a sprint-based approach, breaking the 4-week training into distinct phases.

#### Week 1: Research and Data Engineering

- **Tasks:** Literature review of existing papers on skin analysis (e.g., "DermNet" analysis). Setting up the Python environment. Collecting raw images via web scraping scripts (using BeautifulSoup/Selenium).
- **Deliverable:** A cleaned, labeled dataset of 4,000 images sorted into directory structures.

#### Week 2: Preprocessing and Prototyping

- **Tasks:** Developing the OpenCV pipeline to detect faces and crop them. Implementing statistical normalization. Training a simple "Baseline Model" (a shallow CNN) to establish a performance floor.
- **Deliverable:** A Python script that takes an image and outputs a cropped, normalized array.

#### Week 3: Deep Learning and Transfer Learning

- **Tasks:** Implementing VGG16 with Transfer Learning. Freezing base layers and adding custom dense layers. Running training jobs on the GPU. Tuning hyperparameters like Learning Rate and Batch Size.
- **Deliverable:** A saved model file (.h5) with >85% accuracy.

#### Week 4: Deployment and Reporting

- **Tasks:** wrapping the model in a Flask API. Testing with real-world selfies. Writing this comprehensive project report. Preparing the final presentation.
  - **Deliverable:** The functional API and the final documentation.
-

## Chapter 4: Steps to Achieve Objectives

### 4.1 System Development Life Cycle (SDLC) Methodology

The development of the Nepika module adhered to the **Iterative Waterfall Model**. While the high-level phases (Requirements, Design, Implementation) were sequential, the internal development of the AI model was highly iterative. In Deep Learning, one rarely builds the perfect model on the first try. Instead, it is a cycle of *Hypothesis -> Experiment -> Analysis -> Refinement*.

1. **Requirement Analysis:** Identifying the input/output formats and accuracy targets.
2. **System Design:** selecting the VGG16 architecture and designing the API endpoints.
3. **Implementation (Iterative):**
  - o Cycle 1: Train on small data, check for code errors.
  - o Cycle 2: Train on full data, check for overfitting.
  - o Cycle 3: Apply augmentation and regularization.
4. **Testing:** rigorous evaluation against a hold-out test set.
5. **Deployment:** Local hosting of the API.

### 4.2 Requirement Analysis and Hardware Specification

Deep Learning is computationally expensive. Training a CNN involves performing billions of matrix multiplications. Standard CPUs are inefficient for this parallel workload. Therefore, the hardware setup was a critical component of the requirement analysis.

**Table 2: Hardware and Software Specifications**

Component	Specification	Purpose
<b>Processor (CPU)</b>	Intel Core i7 (10th Gen) or AMD Ryzen 7	Data preprocessing and file I/O operations.
<b>Graphics Card (GPU)</b>	NVIDIA RTX 3060 (12GB VRAM)	Acceleration of Tensor operations (Matrix Multiplication).

<b>RAM</b>	32 GB DDR4	Holding large batches of image data in memory.
<b>Storage</b>	512 GB NVMe SSD	Fast read speeds to feed the GPU without bottlenecks.
<b>OS</b>	Ubuntu 20.04 LTS	Preferred environment for TensorFlow stability.
<b>Python</b>	Version 3.8	Core programming language.
<b>Frameworks</b>	TensorFlow 2.5, Keras	Deep Learning libraries.

### 4.3 Data Collection, Ethics, and Augmentation Strategies

#### Data Sourcing Challenges:

Medical data is usually protected by privacy laws (HIPAA). We could not access patient records. Therefore, we relied on two primary sources:

- Public Research Datasets:** Such as the "Acne04" dataset and subsets of "DermNet NZ". These provided high-quality, clinical-style images.
- Web Scraping:** To simulate real-world usage (selfies), we scraped public images from social media and search engines using keywords like "acne selfie," "dark circles," and "clear skin."

#### Data Cleaning and Labeling:

Scraped data is noisy. It includes memes, drawings, and blurry photos. A manual cleaning process was undertaken to remove irrelevant images. The remaining images were manually labeled into directories: dataset/train/acne, dataset/train/clear, etc.

#### Data Augmentation:

A neural network is only as good as its data. If we only train on faces looking straight ahead, the model will fail if a user tilts their head. To fix this, we used Data Augmentation. This involves generating new training examples from existing ones by applying random transformations.

- Rotation:** Randomly rotating images by  $\pm 20^\circ$ .
- Horizontal Flip:** Mirroring the image (a pimple on the left cheek is structurally the same as one on the right).
- Zoom:** Random zooming to simulate the user holding the camera closer or further away.

- **Brightness Shift:** Simulating different lighting conditions.

This effectively multiplied our dataset size by a factor of 5, providing the model with a richer variety of examples to learn from.

---

*(Placeholder: Visual showing Raw Image -> Augmentation (Flip/Rotate) -> Normalized Tensor.)*

## 4.4 Theoretical Framework: Deep Learning and CNNs

To understand the solution, one must understand the tool: the Convolutional Neural Network (CNN).

The Perceptron vs. CNN:

Traditional Multi-Layer Perceptrons (MLPs) flatten an image into a single long line of pixels. This destroys spatial information (e.g., that pixel A is next to pixel B). CNNs preserve this structure.

The Convolution Operation:

The core mechanism is the "Convolution," a mathematical operation where a small matrix (the Kernel or Filter) slides over the input image.

$$\$$(I * K)(i, j) = \sum_m \sum_n I(i-m, j-n) K(m, n)$$$$

Where \$I\$ is the image and \$K\$ is the kernel.

This operation allows the network to detect "features."

- **Low-Level Features:** Edges, vertical lines, horizontal lines.
- **Mid-Level Features:** Curves, circles, corners (eyes, nose shapes).
- **High-Level Features:** Complex textures (acne lesions, wrinkles).

Activation Function (ReLU):

After every convolution, we apply an activation function. We used the Rectified Linear Unit (ReLU):

$$\$f(x) = \max(0, x)$$$$

ReLU introduces non-linearity. Without it, the entire neural network would collapse into a single linear regression model, incapable of learning complex patterns like skin conditions.

Pooling Layers:

Pooling reduces the size of the image representation, making the computation manageable and

the features "invariant" to small translations. We used Max Pooling, which selects the largest number in a  $2 \times 2$  grid.

---

(Placeholder: A diagram showing a  $3 \times 3$  filter sliding over a  $5 \times 5$  input matrix to produce a feature map.)

## 4.5 Mathematics of Backpropagation and Optimization

How does the machine "learn"? It learns by making mistakes and correcting them.

1. **Forward Pass:** The image goes through the network, and the model guesses "Acne" (Probability 0.7).
  2. Loss Calculation: We compare the guess to the actual label (e.g., "Clear Skin"). We use the Categorical Cross-Entropy Loss function to measure the error:  
$$\text{Loss} = - \sum_{c=1}^M y_{o,c} \log(p_{o,c})$$
  3. **Backpropagation:** We calculate the gradient of the Loss with respect to every weight in the network using the Chain Rule of Calculus. This tells us *direction* we need to adjust the weights to reduce the error.
  4. **Optimization (Adam):** We update the weights. We utilized the **Adam (Adaptive Moment Estimation)** optimizer. Unlike standard Stochastic Gradient Descent (SGD) which uses a fixed learning rate, Adam adapts the learning rate for each parameter, leading to faster convergence.
-

# Chapter 5: Coding and Implementation

## 5.1 Development Environment Setup (Anaconda/CUDA)

The implementation began with setting up a robust development environment. We utilized **Anaconda**, a package manager for data science, to create an isolated virtual environment. This prevents library conflicts (e.g., conflicting versions of NumPy).

### Setup Commands:

Bash

```
conda create -n nepika_env python=3.8
conda activate nepika_env
pip install tensorflow-gpu opencv-python matplotlib flask
```

Crucially, since we were using an NVIDIA GPU, we had to install the **CUDA Toolkit** and **cuDNN (CUDA Deep Neural Network library)**. These libraries allow TensorFlow to interface directly with the GPU hardware, speeding up training by approximately 20x compared to CPU training.

## 5.2 Algorithmic Design and Workflow

The software architecture was designed to be modular. Instead of a single monolithic script, the code was split into functional modules:

1. config.py: Stored variables like Batch Size, Image Dimensions (\$224 \times 224\$), and paths.
2. preprocessing.py: Functions for loading and cleaning images.
3. model.py: The definition of the CNN architecture.
4. train.py: The training loop.
5. app.py: The Flask web server.

## 5.3 Model Architecture: Transfer Learning with VGG16

Training a CNN from scratch requires millions of images and massive compute power. Instead, we used Transfer Learning.

We adopted the VGG16 architecture, pre-trained on the ImageNet dataset (which contains 1.4 million images of general objects).

Why VGG16?

VGG16 is famous for its simplicity. It uses only  $3 \times 3$  convolutional layers stacked on top of each other. This uniformity makes it easy to implement and modify.

**The Modification Strategy:**

1. **Load Base:** We loaded VGG16 but excluded the "top" (the final classification layers that detect dogs/cars).
2. **Freeze Weights:** We set `layer.trainable = False` for the VGG16 layers. This preserves the "knowledge" the model already has about edges and shapes.
3. **Add Custom Head:** We added new layers specifically for our task:
  - o **Global Average Pooling:** Collapses the spatial features into a vector.
  - o **Dense (1024 neurons, ReLU):** A large layer to learn the specific combinations of features that constitute "Acne."
  - o **Dropout (0.5):** Randomly turns off 50% of neurons during training to force the network to be redundant and robust (preventing overfitting).
  - o **Dense (3 neurons, Softmax):** The final output layer for our 3 classes (Acne, Clear, Dark Circles).

---

(Placeholder: Diagram showing the frozen convolutional base and the new trainable dense layers.)

## 5.4 Detailed Code Implementation and Logic

Below is a detailed breakdown of the critical training logic in `train.py`.

Python

```
# Import necessary libraries
from tensorflow.keras.preprocessing.image import ImageDataGenerator
from tensorflow.keras.applications import VGG16
from tensorflow.keras.layers import Dense, GlobalAveragePooling2D, Dropout
```

```

from tensorflow.keras.models import Model
from tensorflow.keras.optimizers import Adam

# 1. Data Generators
# We use generators to load images in batches, so we don't run out of RAM.
train_datagen = ImageDataGenerator(
    rescale=1./255,      # Normalize pixel values to 0-1
    rotation_range=20,   # Data Augmentation
    horizontal_flip=True,
    validation_split=0.2 # Use 20% of data for validation
)

train_generator = train_datagen.flow_from_directory(
    'dataset/train',
    target_size=(224, 224), # VGG16 standard input
    batch_size=32,
    class_mode='categorical',
    subset='training'
)

validation_generator = train_datagen.flow_from_directory(
    'dataset/train',
    target_size=(224, 224),
    batch_size=32,
    class_mode='categorical',
    subset='validation'
)

# 2. Build Model
base_model = VGG16(weights='imagenet', include_top=False, input_shape=(224, 224, 3))

# Freeze the base model
for layer in base_model.layers:
    layer.trainable = False

# Add custom classification head
x = base_model.output
x = GlobalAveragePooling2D()(x)
x = Dense(1024, activation='relu')(x)
x = Dropout(0.5)(x) # Regularization
predictions = Dense(3, activation='softmax')(x) # 3 Classes

model = Model(inputs=base_model.input, outputs=predictions)

```

```

# 3. Compile
model.compile(optimizer=Adam(learning_rate=0.0001),
              loss='categorical_crossentropy',
              metrics=['accuracy'])

# 4. Train
history = model.fit(
    train_generator,
    epochs=25,
    validation_data=validation_generator
)

# 5. Save
model.save('nepika_skin_model.h5')

```

#### Code Logic Explanation:

- **ImageDataGenerator:** This is a crucial Keras utility. It handles the loading, resizing, normalizing, and augmenting of images on the fly. Without this, we would have to write complex loops to process images manually.
- **flow\_from\_directory:** This function infers the class labels based on the folder names. If we have folders acne and clear, it automatically assigns them labels 0 and 1.
- **Adam(learning\_rate=0.0001):** We use a very low learning rate. Since we are using a pre-trained model, we don't want to make drastic changes to the weights; we just want to "nudge" them towards our specific task.

## 5.5 API Development with Flask

The model file (.h5) is useless on its own. It needs an interface. We built a **REST API** using **Flask**.

The API has a single endpoint: /predict (POST).

#### Workflow:

1. Receive image file from request.
2. Read file into memory using `request.files['image'].read()`.
3. Convert the binary stream into a NumPy array using `np.fromstring`.
4. Decode the array into an image using `cv2.imdecode`.
5. **Preprocessing:** This must match the training preprocessing *exactly*.
  - Resize to \$224 \times 224\$.
  - Normalize (divide by 255).
  - Expand dimensions (convert shape from (224, 224, 3) to (1, 224, 224, 3) because the

- model expects a batch).
6. **Inference:** model.predict(image).
  7. **Response:** Format the output as JSON.

### Sample API Code:

Python

```
@app.route('/predict', methods=)
def predict():
    try:
        file = request.files['file']
        #... image processing steps...
        preds = model.predict(img)
        # preds is a list like [[0.02, 0.95, 0.03]]

        class_idx = np.argmax(preds)
        confidence = float(preds[class_idx])

    return jsonify({
        'status': 'success',
        'prediction': classes[class_idx],
        'confidence': confidence
    })
except Exception as e:
    return jsonify({'error': str(e)})
```

---

## 5.6 Testing, Validation, and Performance Metrics

The model was trained for 25 epochs. We monitored two graphs: **Loss** and **Accuracy**.

- **Training Accuracy:** Reached 94%.
- **Validation Accuracy:** Reached 89%.

The slight gap (5%) indicates mild overfitting, but it is within acceptable limits.

Confusion Matrix:

A confusion matrix allows us to see exactly where the model is failing.

**Table 5: Confusion Matrix**

<b>Actual / Predicted</b>	<b>Acne</b>	<b>Clear</b>	<b>Dark Circles</b>
<b>Acne</b>	<b>185</b>	10	5
<b>Clear</b>	12	<b>178</b>	10
<b>Dark Circles</b>	8	15	<b>177</b>

*Analysis:* The model occasionally confuses "Dark Circles" with "Clear Skin," likely because lighting shadows on clear skin can look like dark circles. It is very good at identifying "Acne."

---

# Chapter 6: Conclusions and Recommendations

## 6.1 Summary of Findings

The "Nepika Skin Condition Prediction Module" project successfully met its primary training objectives. By leveraging the VGG16 architecture and Transfer Learning, we constructed a system capable of classifying skin conditions with an accuracy approaching 89%. This confirms the hypothesis that consumer-grade hardware and open-source libraries (TensorFlow/Keras) are sufficient to build high-utility medical triage tools.

The training highlighted that in Applied AI, **Data is King**. The improvements in accuracy came not from making the model more complex, but from cleaning the data and using aggressive Data Augmentation. The project also demonstrated the viability of Python Flask as a lightweight serving layer for AI models, providing a seamless bridge between the complex mathematics of the neural network and the user-friendly interface of a mobile app.

## 6.2 Critical Analysis of Limitations

While functional, the system has specific limitations that must be acknowledged:

1. **Lighting Dependency:** The model is highly sensitive to lighting. Images taken in dim light are often misclassified. A "Low Light Enhancement" preprocessing step is needed.
2. **Skin Tone Bias:** The majority of the training data (from public datasets) features lighter skin tones (Fitzpatrick Types I-III). Consequently, the model's accuracy on darker skin tones (Types IV-VI) is lower. This is a common ethical issue in AI that requires a concerted effort to collect more diverse data.
3. **Static Analysis:** The model analyzes a single frame. It cannot use temporal context (e.g., comparing today's selfie to yesterday's to see if a pimple is healing).

## 6.3 Future Scope: Real-time Analysis and GANs

To evolve Nepika into a market-leading product, the following future works are recommended:

1. **Semantic Segmentation:** Moving beyond classification (saying "You have acne") to segmentation (drawing a polygon around every individual pimple). Architectures like

**U-Net or Mask R-CNN** should be explored for this.

2. **Generative Adversarial Networks (GANs):** Implementing a "Simulation Mode" where the app uses GANs to show the user what their skin would look like *after* treatment, serving as a powerful motivation tool.
  3. **Edge AI (TensorFlow Lite):** Currently, the inference runs on a server. Converting the model to TFLite would allow it to run directly on the phone's processor. This improves privacy (photos never leave the phone) and allows the app to work offline.
-

# Appendices

## Appendix A: Source Code Snippets

### A.1 Requirements.txt

```
tensorflow-gpu==2.5.0
flask==2.0.1
opencv-python==4.5.3
numpy==1.19.5
pillow==8.3.1
```

### A.2 Face Detection Utility

Python

```
# Utilizing Haar Cascades for fast face detection
def crop_face(image_path):
    img = cv2.imread(image_path)
    gray = cv2.cvtColor(img, cv2.COLOR_BGR2GRAY)
    faces = face_cascade.detectMultiScale(gray, 1.3, 5)
    if len(faces) == 0:
        return None
    (x, y, w, h) = faces
    return img[y:y+h, x:x+w]
```